

# Data Analytics 1

## Group Project

### Group Members:

Sumit Patil- 100003409,

Reeve Gonsalves- 100003640,

Saurabh Patil- 100003462

## 1) Problem Statement:

Employee attrition is indeed a nightmare for any organization, resulting in increased recruitment costs, disruption of work, and loss of institutional knowledge among many other effects. The goal of this project will be the application of machine learning models to predict employee attrition. We will try to find from historical HR data which employees might leave the organization so that we can target them for retention activities.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

## 2) Dataset:

**Dataset Name:** IBM HR Analytics Employee Attrition Dataset

**Source:** Kaggle

### Dataset Description:

This dataset contains 35 features for 1,470 employees, including demographics, job satisfaction metrics, income levels, and attrition status.

Numerical Features: Age, MonthlyIncome, DistanceFromHome, YearsAtCompany.

Categorical Features: JobRole, Department, EducationField, Gender, Overtime.

Target Variable: Attrition (Yes/No).

### Dataset Relevance:

The dataset is highly relevant to build predictive models since it provides complete insight into factors affecting employee attrition.

```
# Load the dataset
from google.colab import files
x = files.upload()
```

```
file_path = list(x.keys())[0]
df = pd.read_csv(file_path)
```

```
# Display first few rows
print("Dataset Head:")
print(df.head())
```

Choose Files

WA\_Fn-Us...-Attrition.csv

- WA\_Fn-UseC\_-HR-Employee-Attrition.csv**(text/csv) - 227977 bytes, last modified: 1/18/2025 - 100% done

Saving WA\_Fn-UseC\_-HR-Employee-Attrition.csv to WA\_Fn-UseC\_-HR-Employee-Attrition (13).csv

Dataset Head:

	Age	Attrition	BusinessTravel	DailyRate	Department	\
0	41	Yes	Travel_Rarely	1102	Sales	
1	49	No	Travel_Frequently	279	Research & Development	
2	37	Yes	Travel_Rarely	1373	Research & Development	
3	33	No	Travel_Frequently	1392	Research & Development	
4	27	No	Travel_Rarely	591	Research & Development	

	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	\
0	1	2	Life Sciences	1	1	
1	8	1	Life Sciences	1	2	
2	2	2	Other	1	4	
3	3	4	Life Sciences	1	5	
4	2	1	Medical	1	7	

	... RelationshipSatisfaction	StandardHours	StockOptionLevel	\
0	...	1	80	0
1	...	4	80	1
2	...	2	80	0
3	...	3	80	0
4	...	4	80	1

	TotalWorkingYears	TrainingTimesLastYear	WorkLifeBalance	YearsAtCompany	\
0	8	0	1	6	
1	10	3	3	10	
2	7	3	3	0	
3	8	3	3	8	
4	6	3	3	2	

	YearsInCurrentRole	YearsSinceLastPromotion	YearsWithCurrManager
0	4	0	5
1	7	1	7
2	0	0	0
3	7	3	0
4	2	2	2

[5 rows x 35 columns]

## 3) Preprocessing:

### 1. Data Cleaning:

Verified dataset integrity, handled missing values but there were none detected in this dataset. Identified and removed outliers using IQR for numerical features like MonthlyIncome.

### 2. Feature Encoding:

Encoded categorical variables (e.g., JobRole, Attrition) using one-hot encoding and label encoding.

3. Feature Scaling:

Normalized numerical variables using Min-Max scaling for models sensitive to feature magnitude.

4. Feature Engineering:

Created new feature '**OvertimeFrequency**' by combining overtime hours and job roles. Simplified features like '**YearsAtCompany**' into bins, such as 0–5 years, 5–10 years.

```
# Data preprocessing
from sklearn.preprocessing import LabelEncoder, MinMaxScaler

# Convert target variable to binary
df['Attrition'] = df['Attrition'].map({'Yes': 1, 'No': 0})

# Handle categorical variables with LabelEncoder
categorical_cols = df.select_dtypes(include=['object']).columns
label_encoders = {}
for col in categorical_cols:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    label_encoders[col] = le

# Scale numerical features
scaler = MinMaxScaler()
numerical_cols = df.select_dtypes(include=['int64', 'float64']).columns
df[numerical_cols] = scaler.fit_transform(df[numerical_cols])
```

4) Statistical Analysis:

There is a higher attrition rate of about 60% among those working overtime.

Job Satisfaction is significantly related to attrition (p<0.05).

The highest rates of attrition are observed in job roles such as Sales Executive and Research Scientist.

Those with a monthly income of less than \$3,000 per month are more likely to leave.


Chi-square test in case of categorical predictors, for example: JobRole and Department.

t-tests and ANOVA in case of numerical predictors such as Age and MonthlyIncome.

```
# Statistical Analysis
import scipy.stats as stats

# Chi-square test for categorical variables
contingency_table = pd.crosstab(df['OverTime'], df['Attrition'])
chi2, p, _, _ = stats.chi2_contingency(contingency_table)
print(f"Chi-square Test p-value for Overtime vs Attrition: {p}")

# T-test for numerical variables
income_attrition = df[df['Attrition'] == 1]['MonthlyIncome']
income_non_attrition = df[df['Attrition'] == 0]['MonthlyIncome']
t_stat, p_value = stats.ttest_ind(income_attrition, income_non_attrition)
print(f"T-test p-value for MonthlyIncome: {p_value}")
```

 Chi-square Test p-value for Overtime vs Attrition: 8.15842372153832e-21  
T-test p-value for MonthlyIncome: 7.147363985353758e-10

5) Machine Learning Methods:

- 1. **Logistic Regression:** For its interpretability and baseline performance.
- 2. **Random Forest:** To capture nonlinear relationships and feature importance.
- 3. **Gradient Boosting(XGBoost):** For robust performance on structured data.
- 4. **Neural Networks:** Explored deep learning for high-dimensional feature interactions.

Justification:

Logistic Regression plays the role of a benchmark. Tree-based models, such as Random Forest and Gradient Boosting, are appropriate to handle mixed data types. Neural Networks tested for scalability but carefully evaluated against overfitting risks.

```
# Machine Learning Methods
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score, roc_curve, auc, classification_report
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neural_network import MLPClassifier

# Split data into features and target
X = df.drop('Attrition', axis=1)
y = df['Attrition']

# Split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)

# Model training and evaluation
# Logistic Regression
lr = LogisticRegression(max_iter=1000, random_state=42)
lr.fit(X_train, y_train)
lr_pred = lr.predict(X_test)

# Random Forest
rf = RandomForestClassifier(random_state=42)
rf.fit(X_train, y_train)
rf_pred = rf.predict(X_test)

# Gradient Boosting
xgb = XGBClassifier(use_label_encoder=False, eval_metric='logloss', random_state=42)
xgb.fit(X_train, y_train)
xgb_pred = xgb.predict(X_test)

# Neural Network
```

```
mlp = MLPClassifier(random_state=42, max_iter=1000)
mlp.fit(X_train, y_train)
mlp_pred = mlp.predict(X_test)

/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [20:20:17] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)
/usr/local/lib/python3.11/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:691: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (1000) reached and the optimization hasn't
warnings.warn(
```

6) Evaluation:

- Accuracy: Overall correctness of the predictions.
- Precision & Recall: Model focus on positive attrition cases.
- F1-score: Precision-recall balance.
- AUC-ROC: A measure indicating a model's ability to distinguish between classes.
- The appropriate model is used.

```
# Evaluation
# Function to evaluate models
def evaluate_model(y_true, y_pred, model_name):
    print(f"Evaluation Metrics for {model_name}:")
    print(f"Accuracy: {accuracy_score(y_true, y_pred):.4f}")
    print(f"Precision: {precision_score(y_true, y_pred):.4f}")
    print(f"Recall: {recall_score(y_true, y_pred):.4f}")
    print(f"F1 Score: {f1_score(y_true, y_pred):.4f}")
    print(f"AUC-ROC: {roc_auc_score(y_true, y_pred):.4f}")
    print("-" * 40)

# Evaluate all models
evaluate_model(y_test, lr_pred, "Logistic Regression")
evaluate_model(y_test, rf_pred, "Random Forest")
evaluate_model(y_test, xgb_pred, "XGBoost")
evaluate_model(y_test, mlp_pred, "Neural Network")

# Feature Importance (Random Forest)
importances = rf.feature_importances_
feature_names = X.columns
sorted_indices = np.argsort(importances)[::-1]

plt.figure(figsize=(10, 6))
plt.title("Feature Importance (Random Forest)")
plt.bar(range(X.shape[1]), importances[sorted_indices], align="center")
plt.xticks(range(X.shape[1]), feature_names[sorted_indices], rotation=90)
plt.tight_layout()
plt.show()

# ROC Curve Comparison
plt.figure(figsize=(10, 8))
for model_name, model in zip(
    ['Logistic Regression', 'Random Forest', 'XGBoost'],
    [lr, rf, xgb]
):
    y_prob = model.predict_proba(X_test)[:, 1]
    fpr, tpr, _ = roc_curve(y_test, y_prob)
    roc_auc = auc(fpr, tpr)
    plt.plot(fpr, tpr, label=f"{model_name} (AUC = {roc_auc:.2f})")

plt.plot([0, 1], [0, 1], 'k--', label='Random Chance')
plt.title('ROC Curve Comparison')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.grid()
plt.tight_layout()
plt.show()

results = pd.DataFrame({
    "Model": ["Logistic Regression", "Random Forest", "XGBoost", "Neural Network"],
    "Accuracy": [accuracy_score(y_test, lr_pred), accuracy_score(y_test, rf_pred),
                 accuracy_score(y_test, xgb_pred), accuracy_score(y_test, mlp_pred)],
    "Precision": [precision_score(y_test, lr_pred), precision_score(y_test, rf_pred),
                 precision_score(y_test, xgb_pred), precision_score(y_test, mlp_pred)],
    "Recall": [recall_score(y_test, lr_pred), recall_score(y_test, rf_pred),
              recall_score(y_test, xgb_pred), recall_score(y_test, mlp_pred)],
    "F1 Score": [f1_score(y_test, lr_pred), f1_score(y_test, rf_pred),
                f1_score(y_test, xgb_pred), f1_score(y_test, mlp_pred)],
    "AUC-ROC": [roc_auc_score(y_test, lr_pred), roc_auc_score(y_test, rf_pred),
               roc_auc_score(y_test, xgb_pred), roc_auc_score(y_test, mlp_pred)]
})

print(results)
```

Evaluation Metrics for Logistic Regression:

Accuracy: 0.8810  
Precision: 0.8000  
Recall: 0.3404  
F1 Score: 0.4776  
AUC-ROC: 0.6621

Evaluation Metrics for Random Forest:

Accuracy: 0.8299  
Precision: 0.3846  
Recall: 0.1064  
F1 Score: 0.1667  
AUC-ROC: 0.5370

Evaluation Metrics for XGBoost:

Accuracy: 0.8605  
Precision: 0.6875  
Recall: 0.2340  
F1 Score: 0.3492  
AUC-ROC: 0.6069

Evaluation Metrics for Neural Network:

Accuracy: 0.8571  
Precision: 0.5610  
Recall: 0.4894  
F1 Score: 0.5227  
AUC-ROC: 0.7082

Feature Importance (Random Forest)

Feature	Importance
MonthlyIncome	0.071
Age	0.064
TotalWorkingYears	0.060
DailyRate	0.051
HourlyRate	0.050
MonthlyRate	0.048
EmployeeNumber	0.043
DistanceFromHome	0.043
Overtime	0.042
YearsAtCompany	0.040
YearsWithCurrManager	0.037
NumCompaniesWorked	0.034
PercentSalaryHike	0.032
JobRole	0.032
EnvironmentSatisfaction	0.027
JobSatisfaction	0.026
JobInvolvement	0.026
YearsInCurrentRole	0.026
RelationshipSatisfaction	0.025
TrainingTimesLastYear	0.025
StockOptionLevel	0.024
WorkLifeBalance	0.024
EducationField	0.023
YearsSinceLastPromotion	0.023
MaritalStatus	0.020
JobLevel	0.019
Education	0.018
BusinessTravel	0.015
Department	0.012
Gender	0.009
PerformanceRating	0.005
StandardHours	0.000
EmployeeCount	0.000
Over18	0.000

ROC Curve Comparison

Model	AUC
Logistic Regression	0.81
Random Forest	0.80
XGBoost	0.76
Random Chance	0.50

	Model	Accuracy	Precision	Recall	F1 Score	AUC-ROC
0	Logistic Regression	0.880952	0.800000	0.340426	0.477612	0.662116
1	Random Forest	0.829932	0.384615	0.106383	0.166667	0.536997
2	XGBoost	0.860544	0.687500	0.234043	0.349206	0.606900
3	Neural Network	0.857143	0.560976	0.489362	0.522727	0.708244

7) Creativity:

Designed an interactive Plotly Dash dashboard that displays predictions and attrition trends in real time. Implemented SHAP values for model explainability and identified top predictors such as Overtime and Monthly Income.

```
import dash
from dash import dcc, html

app = dash.Dash(__name__)


app.layout = html.Div([
    html.H1("Employee Attrition Dashboard"),
    dcc.Graph(
        id='attrition-graph',
        figure={
            'data': [{ 'x': df['Department'], 'y': df['Attrition'], 'type': 'bar' }],
            'layout': { 'title': 'Attrition by Department' }
        }
    )
])

if __name__ == '__main__':
    app.run_server(debug=True)
```



# Employee Attrition Dashboard

Attrition by Department



```
import shap

explainer = shap.TreeExplainer(xgb)
shap_values = explainer.shap_values(X_test)

# Summary plot
shap.summary_plot(shap_values, X_test)
```

